

## Install Docker

sudo yum install docker -y

```
[ec2-user@ip-172-31-92-18 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:09:56 ago on Wed Sep 11 15:19:39 2024.
Dependencies resolved.
```

```
=====
Package                                                    Architecture
=====
Installing:
  docker                                                    x86_64
Installing dependencies:
  containerd                                                x86_64
  iptables-libs                                             x86_64
  iptables-nft                                              x86_64
  libcggroup                                                x86_64
  libnetfilter_conntrack                                   x86_64
  libnfnetlink                                              x86_64
  libnftnl                                                  x86_64
  pigz                                                       x86_64
  runc                                                       x86_64

Transaction Summary
```

Then, configure cgroup in a daemon.json file by using following commands

- `cd /etc/docker`
- ```
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
```
- `sudo systemctl enable docker`

- `sudo systemctl daemon-reload`
- `sudo systemctl restart docker`
- `docker -v`

```
[ec2-user@ip-172-31-81-63 docker]$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
docker -v

Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
Docker version 25.0.5, build 5dc9bcc
[ec2-user@ip-172-31-81-63 docker]$
```

#### 4. Install Kubernetes on all 3 machines

SELinux needs to be disabled before configuring kubelet

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee
/etc/apt/sources.list.d/kubernetes.list
```

```
[ec2-user@ip-172-31-81-63 docker]$ sudo setenforce 0
[ec2-user@ip-172-31-81-63 docker]$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

Add kubernetes repository (paste in terminal)

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo[kubernetes]
name=Kubernetes baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/enabled=1
gpgcheck=1 gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cniEOF
```

Type following commands:

- `sudo yum update`

```
sudo yum install -y kubelet kubeadm kubectl  
--disableexcludes=kubernetes
```

```
[ec2-user@ip-172-31-81-63 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes  
Last metadata expiration check: 0:01:34 ago on Wed Sep 11 15:39:05 2024.  
Dependencies resolved.  
=====
```

| Package                  | Architecture | Version                |
|--------------------------|--------------|------------------------|
| Installing:              |              |                        |
| kubeadm                  | x86_64       | 1.30.4-150500.1.1      |
| kubectl                  | x86_64       | 1.30.4-150500.1.1      |
| kubelet                  | x86_64       | 1.30.4-150500.1.1      |
| Installing dependencies: |              |                        |
| conntrack-tools          | x86_64       | 1.4.6-2.amzn2023.0.2   |
| cri-tools                | x86_64       | 1.30.1-150500.1.1      |
| kubernetes-cni           | x86_64       | 1.4.0-150500.1.1       |
| libnetfilter_cthelper    | x86_64       | 1.0.0-21.amzn2023.0.2  |
| libnetfilter_cttimeout   | x86_64       | 1.0.0-19.amzn2023.0.2  |
| libnetfilter_queue       | x86_64       | 1.0.5-2.amzn2023.0.2   |
| socat                    | x86_64       | 1.7.4.2-1.amzn2023.0.2 |

```
Transaction Summary  
=====
```

|         |             |
|---------|-------------|
| Install | 10 Packages |
|---------|-------------|

After installing Kubernetes, we need to configure internetoptions to allow bridging.

```
sudo swapoff -a  
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee  
-a/etc/sysctl.conf  
sudo sysctl -p
```

## 5. Perform this ONLY on the Master machine

Initialize kubernetes by typing below command

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16  
--ignore-preflight-errors=all
```

```
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.81.63:6443 --token zh5jbb.a6ty3eujzc51d15d \
  --discovery-token-ca-cert-hash sha256:0822f656bf52a17a2b6686c123f811306f41495ca650a0aed9bf6cd2d2f6f8c5
[ec2-user@ip-172-31-81-63 docker]$ mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config
[ec2-user@ip-172-31-81-63 docker]$
```

Copy the mkdir and chown commands from the top and execute them

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Copy this join link and save it in clipboard (copy from your output as it differs for each instance)

```
kubeadm join 172.31.81.63:6443 --token zh5jbb.a6ty3eujzc51d15d \
  --discovery-token-ca-cert-hash
sha256:0822f656bf52a17a2b6686c123f811306f41495ca650a0aed9bf6c d2d2f6f8c5
```

Then, add a common networking plugin called flannel file as mentioned in the code.

```
kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

```
[ec2-user@ip-172-31-81-63 docker]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
```

Check the created pod using this command

- kubectl get pods

## 6. Perform this ONLY on the worker machines

Paste the below command on all 2 worker machines

- Sudo yum install iproute-tc -y
  - sudo systemctl enable kubelet
  - sudo systemctl restart kubelet
- 
- kubeadm join 172.31.81.63:6443 --token zh5jbb.a6ty3eujzc51d15d \
- discovery-token-ca-cert-hash  
sha256:0822f656bf52a17a2b6686c123f811306f41495ca650a0aed9bf6cd2d2f6f8  
c5

Now we can see in the master/control node of kubernetes that worker nodes are connected by typing **watch kubectl get nodes** in the **master node instance**

```
Every 2.0s: kubectl get nodes
```

| NAME                          | STATUS | ROLES         | AGE   | VERSION |
|-------------------------------|--------|---------------|-------|---------|
| ip-172-31-81-63.ec2.internal  | Ready  | control-plane | 29m   | v1.30.4 |
| ip-172-31-87-137.ec2.internal | Ready  | <none>        | 5m58s | v1.30.4 |
| ip-172-31-92-18.ec2.internal  | Ready  | <none>        | 5m53s | v1.30.4 |